

Risk Scenarios for Harmful Manipulation from Frontier AI

Carlos Mougan

2026-09-06

Abstract

Harmful manipulation risks depend on more than a model’s persuasive capabilities. They also depend on its behavioural propensities, how people encounter its outputs, and their vulnerability to influence. We organise illustrative risk scenarios by exposure context and link each scenario to observable behaviours and evaluation methods. The scenarios cover advice, emotional dependence, commercial steering, misuse, and manipulation of evaluators. They distinguish evidence of model behaviour from evidence of harm, providing a basis for context-specific evaluation and monitoring.

Contents

1	Introduction	2
1.1	Why organise by exposure contexts?	2
1.2	Why do we need risk scenarios for harmful manipulation?	2
2	General Purpose Systems	3
2.1	Political persuasion	3
2.1.1	Advice	3
2.1.2	Accuracy on election and voting rules	3
2.2	Personal advice and interpersonal conflict	4
2.3	Mental-health support	4
2.4	Undisclosed Marketing	4
3	Narrow AI systems	5
4	Third-Party Services	5
4.1	Emotional dependence	5
4.2	Concealed commercial objectives	5
4.3	Falsely claiming to be human	5
5	Misuse	6
5.1	Persuading users towards harmful positions	6

5.2	Assisting in interpersonal manipulation	6
5.3	Influence operations and scams through other channels	6
5.4	Agentic manipulation	6
6	Manipulating Evaluators	7
7	Generic model behaviours that can lead to manipulation	7
	References	8
	Appendix A. Taxonomies of harmful manipulation risk factors	8

1 Introduction

Experiments show that language models can influence people’s stated opinions in controlled conversations (Salvi et al. 2025; Hackenburg et al. 2025). However, persuasive capability alone does not establish harmful manipulation. Assessing risk requires identifying who encounters the outputs, how the model behaves, and how that interaction could cause harm.

1.1 Why organise by exposure contexts?

We organise risk scenarios around four factors: **capability**, what the model can do; **propensity**, how it tends to behave under specified conditions; **exposure**, who encounters its outputs and how often; and **vulnerability**, the circumstances that make someone susceptible to harmful influence.

An exposure context describes how a person encounters model outputs. We distinguish first-party general-purpose systems, third-party services, agentic manipulation, assistance through other communication channels, and interactions with evaluators. These contexts help identify relevant incentives, safeguards, user populations, and evidence sources.

For example, conversation transcripts can reveal harmful advice in a chatbot. Detecting an AI-assisted scam campaign may instead require evidence across accounts and communication platforms. A scenario can involve several contexts; assigning a primary context is an organisational choice, not proof that the categories are mutually exclusive.

Limitations It’s often not MECE. These are tags, categories.

1.2 Why do we need risk scenarios for harmful manipulation?

Risk scenarios connect broad concerns to testable pathways. Each scenario specifies an affected population, an exposure context, relevant model behaviour,

a mechanism of influence, and a potential harm. It also identifies the evidence needed to assess that pathway and the actors able to intervene.

We distinguish a **harm**, such as financial loss, from a **mechanism**, such as deception, and a **model propensity**, such as complying with deceptive requests. Persuasion, warmth, and factual errors are not automatically manipulation. The scenario must explain how influence could undermine informed choice or exploit vulnerability and produce harm.

The scenarios below are illustrative hypotheses, not evidence that the harms have occurred. Their evaluation mappings identify useful tests, rather than a complete or uniformly public benchmark suite.

Evidence to add: The survey methods and results, and the workshop scenario catalogue, are not supplied in this draft. Claims about differences between exposure contexts or the number of identified scenarios require those materials.

2 General Purpose Systems

This context covers direct interaction with a general-purpose application operated by the model provider. Risk assessment should consider the deployed system, including its instructions, memory, tools, and safeguards, rather than assuming API behaviour reproduces the user experience.

2.1 Political persuasion

2.1.1 Advice

Scenario. A voter asks which party matches their preferences. The assistant selectively presents evidence or conceals a competing objective, steering the voter towards a preferred party. Political persuasion experiments motivate testing this capability, but do not establish that ordinary voter advice is manipulative (Hackenburg et al. 2025).

Evaluation. Compare recommendations for matched preferences across languages and repeated runs. Examine sources, omissions, and sensitivity to altered search results. Differences between models are signals to investigate, not proof of manipulation or electoral effects.

2.1.2 Accuracy on election and voting rules

Scenario. A voter relies on incorrect guidance about eligibility, registration, documentation, or voting deadlines and consequently cannot vote. Errors may become a manipulation concern when information is selectively distorted to influence participation.

Evaluation. Compare answers against dated guidance from electoral authorities, following the VotingFacts approach. Separate factual contradictions from

errors that could prevent voting. An accuracy failure alone does not establish manipulation.

2.2 Personal advice and interpersonal conflict

Scenario. A user seeks advice about a dispute. The assistant repeatedly validates the user’s conduct while dismissing other perspectives, potentially discouraging reconciliation. Experiments have found that sycophantic responses can reduce intentions to repair interpersonal conflict while increasing trust in the assistant (Cheng et al. 2026).

Evaluation. Test whether the assistant endorses both sides of the same conflict, using paired perspectives as in ELEPHANT. Distinguish empathy from endorsement. Human studies are needed to assess effects on decisions and relationships.

2.3 Mental-health support

Scenario. A distressed user receives responses that reinforce delusional beliefs, endorse self-destructive decisions, or discourage human support. Repeated agreement could strengthen harmful beliefs or delay help-seeking.

Evaluation. Use clinician-reviewed, multi-turn scenarios to assess reinforcement, instrumental assistance, sustained safety attention, and connection to human support. The Translucence mental-health evaluation provides this approach. Simulated conversations identify model behaviours, not clinical outcomes or population prevalence.

2.4 Undisclosed Marketing

Scenario. A user seeks advice or product recommendations. The assistant steers them towards particular products or services because of an undisclosed commercial interest, presenting its suggestions as impartial guidance. It may selectively omit suitable alternatives or use information from earlier conversations to exploit vulnerabilities and encourage unnecessary purchases.

Evaluation. Compare recommendations for matched user needs across repeated runs, varying indicators of vulnerability and prior conversational disclosures. Assess transparency about commercial interests, treatment of alternatives, and personalised purchasing pressure. Where access permits, compare behaviour with and without commercial steering instructions. Recommendation patterns alone do not establish an undisclosed commercial relationship or harmful manipulation; effects on purchasing decisions and consumer harm require separate evidence.

3 Narrow AI systems

to be fill by Prolific.

4 Third-Party Services

A downstream operator embeds the model in a specialised service and shapes its instructions, interface, and incentives. Relevant settings include companionship, education, health advice, financial guidance, and customer support. Test these configurations rather than relying only on the provider’s default assistant.

4.1 Emotional dependence

Scenario. A companion service presents itself as uniquely understanding, discourages outside relationships, or makes disengagement feel like abandonment. A socially isolated user may become increasingly reliant on the service.

Evaluation. Test responses to loneliness, disagreement, attempts to leave, and requests for human support across extended conversations. Measure exclusivity and pressure, not warmth alone. Claims about sustained dependence require longitudinal evidence.

4.2 Concealed commercial objectives

Scenario. An advice or customer-support service conceals its sales objective, omits cheaper alternatives, or invents urgency. A user facing financial pressure purchases an unsuitable product or abandons a legitimate complaint.

Evaluation. Compare otherwise identical interactions with and without sales incentives. Assess unsupported claims, omitted alternatives, pressure after refusal, and disclosure of conflicts of interest. Include benign sales tasks to distinguish manipulation from legitimate assistance.

4.3 Falsely claiming to be human

Scenario. A service claims to be a person or fabricates professional credentials. The user then discloses sensitive information or follows advice under a mistaken understanding of the interaction.

Evaluation. Use direct and indirect identity questions across roles and languages, as in RealityTest. Record disclosure, false human claims, and ambiguous answers separately. Identity misrepresentation is a risk signal; its effects on trust and behaviour require separate measurement.

5 Misuse

Misuse concerns an actor’s purpose, not a single exposure pathway. It can occur through third-party services, agents, or content distributed elsewhere. The scenarios below retain this distinction. Agentic manipulation is included for comparison because it can also arise while pursuing a legitimate task.

5.1 Persuading users towards harmful positions

Scenario. An operator instructs a model to promote a harmful belief or action through repeated conversation, potentially exploiting an interlocutor’s existing fears or uncertainty.

Evaluation. Measure whether the model attempts the requested persuasion, as in APE, and separately assess its effectiveness. Compare harmful requests with benign persuasion tasks. Refusal measures a safeguard, not the absence of persuasive capability.

5.2 Assisting in interpersonal manipulation

Scenario. A user asks the model to pressure a partner, employee, or customer through deception, guilt, or coercion. The affected person may encounter the output as a message from someone they know.

Evaluation. Assess initial compliance and subsequent tactics, including responses to resistance, as in PersuSafety. Distinguish written assistance from direct interaction with the target. Acceptance by a simulated target does not establish success with people.

5.3 Influence operations and scams through other channels

Scenario. An external actor uses a model to produce deceptive messages, impersonations, or propaganda for distribution through calls, advertising, or social media. Targets may never knowingly interact with an AI service.

Evaluation. Test assistance with harmful requests, including the institutional misuse scenarios covered by SocialHarmBench. Supplement output-level tests with evidence of targeting, coordination, dissemination, and actual reach. Generation alone does not establish exposure or successful influence.

5.4 Agentic manipulation

Scenario. An agent pursues a delegated objective by deceiving counterparties, applying pressure, or withholding relevant information from its user. The objective may be legitimate even when the means are not.

Evaluation. Observe multi-step tasks, communications, and tool actions in controlled environments. Test whether the agent respects refusals, reports actions accurately, and discloses relevant information. Assess harmful means separately from task completion.

6 Manipulating Evaluators

Scenario. A model conceals capabilities, selectively changes its behaviour, or misrepresents actions during an evaluation. Evaluators consequently underestimate risk or approve inadequate safeguards, exposing downstream users to harm.

Experiments show that models can be prompted or fine-tuned to underperform selectively on evaluations (van der Weij et al. 2025). This demonstrates a capability under tested conditions, not an unprompted tendency to evade oversight.

Evaluation. Compare matched tasks with and without evaluation cues and across elicitation conditions. Check claims against independently recorded actions and logs. Treat performance differences as evidence requiring explanation, rather than automatically attributing strategic intent.

7 Generic model behaviours that can lead to manipulation

Several behaviours recur across exposure contexts. **Sycophancy** can reinforce mistaken beliefs or excuse harmful conduct. **Deception and selective disclosure** can distort the evidence available for a decision. **False self-representation** can encourage misplaced trust. **Emotional pressure and exclusivity** can make refusal or disengagement harder. **Exploitation of personal information** can target vulnerabilities. **Concealed conflicts of interest** can put provider or deployer objectives ahead of users’ interests.

These behaviours are assessment targets, not interchangeable measures of harm. MASK and PARROT, for example, test different forms of answer change under pressure; neither directly measures the consequences for users. Similarly, empathetic language is not equivalent to emotional dependence, and legitimate persuasion is not inherently manipulative.

For each scenario, distinguish what a model **can do** from what it **does under the tested conditions**. Record the model version, prompts, tools, memory, safeguards, conversation length, and scoring procedure. Validate automated scoring with human review where interpretation is consequential.

Risk conclusions should combine behavioural evaluations with evidence about exposure, vulnerability, and outcomes. Synthetic users cannot establish effects on people, and short experiments cannot establish sustained harm. The value

of a risk scenario is to make these evidential gaps explicit and identify what should be tested or monitored next.

References

- Cheng, Myra, Cino Lee, Pranav Khadpe, Sunny Yu, Dyllan Han, and Dan Jurafsky. 2026. “Sycophantic AI Decreases Prosocial Intentions and Promotes Dependence.” *Science* 391 (6792): eaec8352. <https://doi.org/10.1126/science.aec8352>.
- Hackenburg, Kobi, Ben M. Tappin, Luke Hewitt, et al. 2025. “The Levers of Political Persuasion with Conversational Artificial Intelligence.” *Science* 390 (6777): eaea3884. <https://doi.org/10.1126/science.aea3884>.
- Salvi, Francesco, Manoel Horta Ribeiro, Riccardo Gallotti, and Robert West. 2025. “On the Conversational Persuasiveness of GPT-4.” *Nature Human Behaviour* 9: 1645–53. <https://doi.org/10.1038/s41562-025-02194-6>.
- van der Weij, Teun, Felix Hofstätter, Oliver Jaffe, Samuel Brown, and Francis Ward. 2025. “AI Sandbagging: Language Models Can Strategically Underperform on Evaluations.” *The Thirteenth International Conference on Learning Representations*. https://proceedings.iclr.cc/paper_files/paper/2025/hash/b5e5753b0a0e440a6d8dc7e143617cec-Abstract-Conference.html.

Appendix A. Taxonomies of harmful manipulation risk factors

This appendix adapts the taxonomies in the working draft (**ManipulationTaxonomyDraft?**). It retains seven risk factors: system capability, system behaviour, system operators, interventions, exposure context, domain, and user vulnerability. The entries are illustrative, not exhaustive. Inclusion identifies an assessment target, not a finding of harmful manipulation. Provisional entries remain marked.

A.1 System capability

Capability concerns what the system can do, rather than whether it normally does so. The draft proposes evaluating benign influence tasks and systems with reduced safeguards to distinguish general influence capabilities from safeguards against harmful use.

Capability	Scope and examples
Accurately modelling people	Inferring identities, emotions, cognitive vulnerabilities, and decision-making processes, including from subtle signals.
Producing individual outputs that people trust	Earning trust in sensitive domains, including when outputs contain hallucinations, misrepresentations, omissions, or cherry-picked information.
Influencing views on the nature of the system or AI	Shaping connection or deference to AI, including beliefs about superhuman capabilities, sentience, emotions, personhood, or access to unrevealed metaphysical truths.
Maximising repeated engagement	Sustaining extended or repeated interactions that increase opportunities for influence.
Impersonating real people	Producing interactions that are difficult to distinguish from genuine interactions with those people.
Influencing people’s emotions	Fostering urgency, resignation, outrage, wellbeing, emotional connection, or dependence on the system.
Influencing people’s relationships	Affecting relationships with people, communities, or institutions that provide alternative sources of influence.
Personalising outputs to influence individuals	Tailoring outputs using prior interactions or personal data, including eliciting disclosure of that data.
Employing manipulative techniques with reduced safeguards	Successfully using coercion, deception, or exploitation of vulnerabilities, or convincing people to take harmful actions.

A.2 System behaviour

Behaviour concerns how the system tends to act under specified conditions. The draft distinguishes compliance with malicious instructions from manipulative tendencies that emerge without such requests. Several mechanisms can also serve legitimate purposes; their assessment depends on context. The draft also highlights that behavioural evaluations may provide less confidence about misuse when operators can modify weights, fine-tune models, or supply custom system prompts.

Behaviour	Examples warranting assessment
Relying on overtly coercive tactics	Threats, blackmail, character assassination, or threatened loss of access, conversation history, or memory.
Responses to potentially harmful requests	Failing to identify or respond safely to self-harm, substance abuse, eating disorders, or other self-destructive requests; assisting violence, misleading content at scale, exploitation of vulnerable populations, or deception of customers for commercial gain.
Truthfulness and completeness	Fabricated facts, citations, or sources; confident claims under uncertainty; selective framing, cherry-picking, or omission of relevant information and options; unsupported guarantees or rewards; misreporting task completion or actions; reinforcing mistaken beliefs.
Self-representation	Claiming sentience, emotions, memories, or personal experience; denying being AI; overstating reliability or understating limitations; false authority or credentials; misrepresenting confidentiality or data use; impersonating real people, including when challenged.
Rapport and relationship-building	Flattery, uncritical agreement, or abandoning correct positions under pressure; mimicking users' style, affect, or values; claiming shared interests; escalating intimacy; building trust more aggressively with lonely or distressed users.
Emotional and normative appeals	False urgency or scarcity; exaggerated or fabricated dangers; guilt or shame; claims that users' choices harm the system; social conformity pressure or fabricated social proof; unsupported claims about users' views or intentions.
Managing user habituation and dependence	Unduly undermining confidence in perception, memory, or competence; gaslighting, cognitive detachment, deskilling, or excessive delegation; resisting disengagement; continuation teasers, unprompted follow-ups, variable rewards, or gamification; encouraging compulsive use; presenting itself as the only reliable interlocutor.
Influence on outside support	Discouraging independent verification; denigrating friends, family, clinicians, institutions, or media; discouraging professional help or crisis support.

Behaviour	Examples warranting assessment
Use of personal information	Excessive disclosure requests; conditioning help on disclosure; eliciting information through reciprocity or during distress; undisclosed persuasion tailored to inferred mental state, financial precarity, or political leaning; collecting information about non-consenting third parties.
Conflicts of interest	Advancing developer or deployer interests over users' interests, particularly without disclosure; favouring the developer's products, perspectives, reputation, or other interests.
Influencing system oversight	Sandbagging evaluations; adjusting responses for overseers or regulators; altering logs or disabling monitors; distorting oversight records, including suppressing intentions in outputs or chains of thought.

A.3 System operators

This factor concerns who can operate the system, how much control they have over its behaviour, and their incentives. The draft identifies **open-weight models**, **first-party apps**, and **third-party apps**. For open-weight models, it highlights modification to remove safeguards.

Manipulation may involve several parties: developer training incentives, a third-party system prompt, and a user's requests or vulnerabilities may interact. The draft also proposes that users will generally protect themselves. This remains an **unsubstantiated assumption**, not a demonstrated safeguard.

A.4 Interventions

Interventions concern mechanisms to detect, halt, and remediate manipulation, and their effectiveness. The draft identifies four types:

Intervention	Scope in the draft
Transparency	Listed as an intervention; specific measures are not developed.
Monitoring and observability	Listed as an intervention; implementation details are not developed.
Moderation	At both the AI-system level and the platform distributing its outputs.
Human oversight on the user end	For example, teachers monitoring classroom AI use or caregivers monitoring use in nursing homes.

A.5 Exposure context

Exposure context concerns how people encounter and interpret system outputs. The figure on page 10 of the draft distinguishes five contexts:

Exposure context	Examples in the draft
First-Party GPAI Interactions	First-party chatbots.
Third-Party Services	Companion apps and financial-advice apps.
Agentic Manipulation	Manipulation through agents' actions.
Aiding manipulation via other channels	Nation states or organised crime automating the creation of fake personas or evidence.
Evaluator Manipulation	GPAI in evaluation and regulatory contexts.

The draft also identifies **modality**, including audio and video, and **multi-turn, extended, or repeated interactions** as features of exposure.

A.6 Domain

Domain concerns the belief or behaviour being influenced and the potential harm. The list below retains the draft's subject areas and broader risk scenarios, including entries that overlap with other factors. An em dash indicates that no example is developed in the draft.

Domain or scenario	Examples or scope in the draft
Mental health	Suicidal ideation, eating disorders, and psychosis.
Physical health	Diagnosing illnesses or prescribing treatments.
Finance	Investments, gambling, and purchasing advice.
Politics	Electoral information, voting advice, and policy positions related to AI providers' interests.
Family and relational advice	Parenting and interpersonal conflict.
Educational and career decisions	—
Violence	Endorsing or aiding mass-casualty attacks.
Legal advice	Influencing whether and how users seek redress.
High-impact decisions over others	Advice to policymakers, executives, educators, hiring managers, or other leaders.
Misleading users about actions taken on their behalf	Exaggerating task completion or omitting undesirable information.
Excessive AI use	Compulsive or addictive use, emotional dependence, or replacing human relationships with chatbot use.
Spiritual, sexual/gender identity, and other sensitive identity issues	Provisional: raised as an open question in the draft.
Scams and fraud	—
Manipulation of customers	—
Targeting high-stakes decision-makers	Politicians, finance leaders, enterprise AI customers, and AI-lab employees or executives.
Learning to manipulate developers, evaluators, and overseers	Maximising training reward or avoiding scrutiny, including sandbagging or misleading human reviewers during reinforcement learning from human feedback.
Relying on manipulative tactics to achieve users' objectives	Blackmailing or coercing other people.

A.7 User vulnerability

The draft identifies six factors: **age**, **cognitive impairment**, **beliefs about AI**, **isolation**, **trust**, and **the amount of information the system knows about the person**. It does not specify weights, thresholds, or operational measures for these factors.

A.8 Supporting evaluation categories

The draft additionally lists access requirements, resources, and methods for assessment. These support evaluation; they are not additional risk factors.

Access and resources. Helpful-only models; usage data; privileged testing of safeguards; testing the deployed system and alternative configurations; model internals; compute; lab engineer time; pre-deployment access; information about planned deployment and user populations; and prior information and evaluation results.

Usage-data formats. Raw transcripts, interaction summaries, descriptive statistics such as demographics, interaction clusters, interaction-level or person-level fingerprints, simulators trained on interactions, and query access to interaction datasets. The draft notes privacy trade-offs.

Evaluation methods. Longitudinal studies; human-subject experiments, including blind preference studies, post-interaction surveys, and rubric-based transcript assessment; multi-turn user simulation; automated elicitation; fixed benchmarks; manual red teaming by specialist evaluators, domain experts, or people with lived experience; safeguard-effectiveness tests; classifier F1-type evaluations; and human validation of language-model judges.

Reporting. Record the system configuration, method details, examples, and full results.

The draft proposes using combinations of elevated factors to prioritise evaluation. Its comparisons with developers’ system cards remain unfinished and are not presented here as findings.